

Image Preprocessing

In [10]:

```
import os, types
import pandas as pd
from botocore.client import Config
import ibm_boto3

def __iter__(self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your c
redentials.
# You might want to remove those credentials before you share the notebook.
cos_client = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='tGcNDmAc_N8W25Ld_PCWuPU19MOPW3PLmzlOU0XjZCI2',
    ibm_auth_endpoint="https://iam.cloud.ibm.com/oidc/token",
    config=Config(signature_version='oauth'),
    endpoint_url='https://s3.private.us.cloud-object-storage.appdomain.cloud')

bucket = 'traincnnmodel-donotdelete-pr-y2q0cidugru0d7'
object_key = 'conversation_engine_for_deaf_and_dumb.zip'

streaming_body_1 = cos_client.get_object(Bucket=bucket, Key=object_key)['Body']

# Your data file was loaded into a botocore.response.StreamingBody object.
# Please read the documentation of ibm_boto3 and pandas to learn more about the possibili
ties to load the data.
# ibm_boto3 documentation: https://ibm.github.io/ibm-cos-sdk-python/
# pandas documentation: http://pandas.pydata.org/
from io import BytesIO
import zipfile
zip_ref = zipfile.ZipFile(BytesIO(streaming_body_1.read()), 'r')
file_paths=zip_ref.namelist()
for path in file_paths:
    zip_ref.extract(path)
```

In [14]:

```
f=os.listdir('/home/wsuser/work/Dataset/training_set')
f
```

Out[14]:

```
['B', 'E', 'F', 'G', 'D', 'I', 'C', 'H', 'A']
```

In [15]:

```
#image Preprocessing
from keras.preprocessing.image import ImageDataGenerator
from keras.preprocessing.image import ImageDataGenerator
train_datagen=ImageDataGenerator(rescale = 1./255, shear_range=0.2, zoom_range=0.2, horiz
ontal_flip=True)
test_datagen=ImageDataGenerator (rescale = 1./255)
x_train = train_datagen.flow_from_directory('/home/wsuser/work/Dataset/training_set', tar
get_size=(64,64), batch_size=300, class_mode='categorical', color_mode = "grayscale")
x_test = train_datagen.flow_from_directory('/home/wsuser/work/Dataset/test_set', target_s
ize=(64,64), batch_size=300, class_mode='categorical', color_mode = "grayscale")
```

Found 15750 images belonging to 9 classes.

Found 2250 images belonging to 9 classes.

Model Building

In [16]:

```

from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Convolution2D
from keras.layers import MaxPooling2D
from keras.layers import Dropout
from keras.layers import Flatten
model=Sequential()
model.add(Convolution2D(32, (3,3), input_shape=(64,64,1), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Flatten())
model.add(Dense(units=512, activation='relu'))
model.add(Dense(units=9, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit_generator(x_train, steps_per_epoch=24, epochs=10, validation_data=x_test, validation_steps=40)

```

/tmp/wsuser/ipykernel_210/3767310412.py:14: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

```

model.fit_generator(x_train, steps_per_epoch=24, epochs=10, validation_data=x_test, validation_steps=40)

```

Epoch 1/10

24/24 [=====] - ETA: 0s - loss: 1.2163 - accuracy: 0.6738WARNING:tensorflow:Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches (in this case, 40 batches). You may need to use the repeat() function when building your dataset.

24/24 [=====] - 23s 950ms/step - loss: 1.2163 - accuracy: 0.6738 - val_loss: 0.5050 - val_accuracy: 0.8818

Epoch 2/10

24/24 [=====] - 20s 846ms/step - loss: 0.2580 - accuracy: 0.9283

Epoch 3/10

24/24 [=====] - 21s 853ms/step - loss: 0.1422 - accuracy: 0.9651

Epoch 4/10

24/24 [=====] - 21s 851ms/step - loss: 0.0847 - accuracy: 0.9803

Epoch 5/10

24/24 [=====] - 21s 850ms/step - loss: 0.0531 - accuracy: 0.9872

Epoch 6/10

24/24 [=====] - 20s 826ms/step - loss: 0.0505 - accuracy: 0.9872

Epoch 7/10

24/24 [=====] - 20s 841ms/step - loss: 0.0361 - accuracy: 0.9926

Epoch 8/10

24/24 [=====] - 21s 863ms/step - loss: 0.0260 - accuracy: 0.9946

Epoch 9/10

24/24 [=====] - 20s 833ms/step - loss: 0.0213 - accuracy: 0.9955

Epoch 10/10

24/24 [=====] - 20s 837ms/step - loss: 0.0188 - accuracy: 0.9959

Out[16]:

<keras.callbacks.History at 0x7f478d08ce20>

In [17]:

```

model.save('cbot.h5')

```

In [19]:

```

tar -zcvf cbotModel.tgz cbot.h5

```

cbot.h5

In [21]:

```

ls

```

cbot.h5 cbotModel.tgz Dataset/

In [22]:

```

pip install watson-machine-learning-client --upgrade

```

Collecting watson-machine-learning-client

In [30]:

```
client.set.default_space(space_uid)
```

Out[30]:

'SUCCESS'

In [31]:

```
client.software_specifications.list()
```

| NAME | ASSET_ID | TYPE |
|-------------------------------|---------------------------------------|------|
| default_py3.6 | 0062b8c9-8b7d-44a0-a9b9-46c416adcdbd9 | base |
| kernel-spark3.2-scala2.12 | 020d69ce-7ac1-5e68-ac1a-31189867356a | base |
| pytorch-onnx_1.3-py3.7-edt | 069ea134-3346-5748-b513-49120e15d288 | base |
| scikit-learn_0.20-py3.6 | 09c5a1d0-9c1e-4473-a344-eb7b665ff687 | base |
| spark-mllib_3.0-scala_2.12 | 09f4cff0-90a7-5899-b9ed-1ef348aebdee | base |
| pytorch-onnx_rt22.1-py3.9 | 0b848dd4-e681-5599-be41-b5f6fccc6471 | base |
| ai-function_0.1-py3.6 | 0cdb0f1e-5376-4f4d-92dd-da3b69aa9bda | base |
| shiny-r3.6 | 0e6e79df-875e-4f24-8ae9-62dcc2148306 | base |
| tensorflow_2.4-py3.7-horovod | 1092590a-307d-563d-9b62-4eb7d64b3f22 | base |
| pytorch_1.1-py3.6 | 10ac12d6-6b30-4ccd-8392-3e922c096a92 | base |
| tensorflow_1.15-py3.6-ddl | 111e41b3-de2d-5422-a4d6-bf776828c4b7 | base |
| autoai-kb_rt22.2-py3.10 | 125b6d9a-5b1f-5e8d-972a-b251688ccf40 | base |
| runtime-22.1-py3.9 | 12b83a17-24d8-5082-900f-0ab31fbfd3cb | base |
| scikit-learn_0.22-py3.6 | 154010fa-5b3b-4ac1-82af-4d5ee5abbc85 | base |
| default_r3.6 | 1b70aec3-ab34-4b87-8aa0-a4a3c8296a36 | base |
| pytorch-onnx_1.3-py3.6 | 1bc6029a-cc97-56da-b8e0-39c3880dbbe7 | base |
| kernel-spark3.3-r3.6 | 1c9e5454-f216-59dd-a20e-474a5cdf5988 | base |
| pytorch-onnx_rt22.1-py3.9-edt | 1d362186-7ad5-5b59-8b6c-9d0880bde37f | base |
| tensorflow_2.1-py3.6 | 1eb25b84-d6ed-5dde-b6a5-3fbdf1665666 | base |
| spark-mllib_3.2 | 20047f72-0a98-58c7-9ff5-a77b012eb8f5 | base |
| tensorflow_2.4-py3.8-horovod | 217c16f6-178f-56bf-824a-b19f20564c49 | base |
| runtime-22.1-py3.9-cuda | 26215f05-08c3-5a41-a1b0-da66306ce658 | base |
| do_py3.8 | 295addb5-9ef9-547e-9bf4-92ae3563e720 | base |
| autoai-ts_3.8-py3.8 | 2aa0c932-798f-5ae9-abd6-15e0c2402fb5 | base |
| tensorflow_1.15-py3.6 | 2b73a275-7cbf-420b-a912-eae7f436e0bc | base |
| kernel-spark3.3-py3.9 | 2b7961e2-e3b1-5a8c-a491-482c8368839a | base |
| pytorch_1.2-py3.6 | 2c8ef57d-2687-4b7d-acce-01f94976dac1 | base |
| spark-mllib_2.3 | 2e51f700-bca0-4b0d-88dc-5c6791338875 | base |
| pytorch-onnx_1.1-py3.6-edt | 32983cea-3f32-4400-8965-dde874a8d67e | base |
| spark-mllib_3.0-py37 | 36507ebe-8770-55ba-ab2a-eafe787600e9 | base |
| spark-mllib_2.4 | 390d21f8-e58b-4fac-9c55-d7ceda621326 | base |
| autoai-ts_rt22.2-py3.10 | 396b2e83-0953-5b86-9a55-7ce1628a406f | base |
| xgboost_0.82-py3.6 | 39e31acd-5f30-41dc-ae44-60233c80306e | base |
| pytorch-onnx_1.2-py3.6-edt | 40589d0e-7019-4e28-8daa-fb03b6f4fe12 | base |
| pytorch-onnx_rt22.2-py3.10 | 40e73f55-783a-5535-b3fa-0c8b94291431 | base |
| default_r36py38 | 41c247d3-45f8-5a71-b065-8580229facf0 | base |
| autoai-ts_rt22.1-py3.9 | 4269d26e-07ba-5d40-8f66-2d495b0c71f7 | base |
| autoai-obm_3.0 | 42b92e18-d9ab-567f-988a-4240ba1ed5f7 | base |
| pmml-3.0_4.3 | 493bcb95-16f1-5bc5-bee8-81b8af80e9c7 | base |
| spark-mllib_2.4-r_3.6 | 49403dff-92e9-4c87-a3d7-a42d0021c095 | base |
| xgboost_0.90-py3.6 | 4ff8d6c2-1343-4c18-85e1-689c965304d3 | base |
| pytorch-onnx_1.1-py3.6 | 50f95b2a-bc16-43bb-bc94-b0bed208c60b | base |
| autoai-ts_3.9-py3.8 | 52c57136-80fa-572e-8728-a5e7cbb42cde | base |
| spark-mllib_2.4-scala_2.11 | 55a70f99-7320-4be5-9fb9-9edb5a443af5 | base |
| spark-mllib_3.0 | 5c1b0ca2-4977-5c2e-9439-ffd44ea8ffe9 | base |
| autoai-obm_2.0 | 5c2e37fa-80b8-5e77-840f-d912469614ee | base |
| spss-modeler_18.1 | 5c3cad7e-507f-4b2a-a9a3-ab53a21dee8b | base |
| cuda-py3.8 | 5d3232bf-c86b-5df4-a2cd-7bb870a1cd4e | base |
| autoai-kb_3.1-py3.7 | 632d4b22-10aa-5180-88f0-f52dfb6444d7 | base |
| pytorch-onnx_1.7-py3.8 | 634d3cdc-b562-5bf9-a2d4-ea90a478456b | base |

Note: Only first 50 records were displayed. To display more use 'limit' parameter.

In [40]:

```
software_spec_uid=client.software_specifications.get_uid_by_name('tensorflow_rt22.1-py3.9')
)
```

```
software_spec_uid
```

Out[40]:

```
'acd9c798-6974-5d2f-a657-ce06e986df4d'
```

In [49]:

```
model_details=client.repository.store_model(model='cbotModel.tgz',meta_props={
    client.repository.ModelMetaNames.NAME:"CNN",
    client.repository.ModelMetaNames.TYPE:"tensorflow_2.7",
    client.repository.ModelMetaNames.SOFTWARE_SPEC_UID:software_spec_uid

})
model_id=client.repository.get_model_id(model_details)
model_id
```

Out[49]:

```
'5ca9eb20-97a7-4bf5-a583-c152ee8ef9ad'
```

In [50]:

```
client.repository.download(model_id,"cbotModel.tar.gz")
```

Successfully saved model content to file: 'cbotModel.tar.gz'

Out[50]:

```
'/home/wsuser/work/cbotModel.tar.gz'
```

Test The Model

In [63]:

```
from keras.models import load_model
import numpy as np
import cv2
from keras.preprocessing import image
```

In [52]:

```
model=load_model('cbot.h5')
```

In [57]:

```
from skimage.transform import resize
def detect(frame):
    img=image.img_to_array(frame)
    img = resize(img,(64,64,1))
    img = np.expand_dims(img,axis=0)
    pred=np.argmax(model.predict(img))
    op=['A','B','C','D','E','F','G','H','I']
    print("THE PREDICTED LETTER IS ",op[pred])
```

In [65]:

```
img=image.load_img("/home/wsuser/work/Dataset/test_set/D/102.png")
detect(img)
```

```
THE PREDICTED LETTER IS  D
```

In []: